



# Learning state representations for formal verification

David Kerkkamp



# Outline

- Motivation and goal
- State representation learning
- Pipeline
- Pong demo
- Conclusion



# Motivation and goal


## Motivation:

- Verified guarantees against dangerous or fatal situations in systems involving AI/ML
- To apply formal verification, a suitable model of the environment is needed


## Goal:

- Learning models of environment suitable for model checking
- Focus on learning state representations of Atari 2600 games

# Atari 2600



Enemy near



Enemy

Falling off edge



# Plan

Create a pipeline that learns state representations and generates a model:

- State representation learning: based on codebase from paper by Anand et al. <sup>[1]</sup>
  - Contains techniques for learning representations of Atari games
- Generate a model: Markov Decision Process

[1] A. Anand, E. Racah, S. Ozair, Y. Bengio, M.-A. Côté, and R. D. Hjelm, “Unsupervised state representation learning in atari,” *arXiv preprint arXiv:1906.08226*, 2019.


# State representation learning

Atari Annotated RAM Interface (AtariARI)

Using OpenAI Gym

Learning representations in two parts:

- Training an encoder
- Predict state labels with linear probing
  - For each label, a 256-way linear classifier is trained



Different encoder architectures are possible



# Codebase

Existing code for:

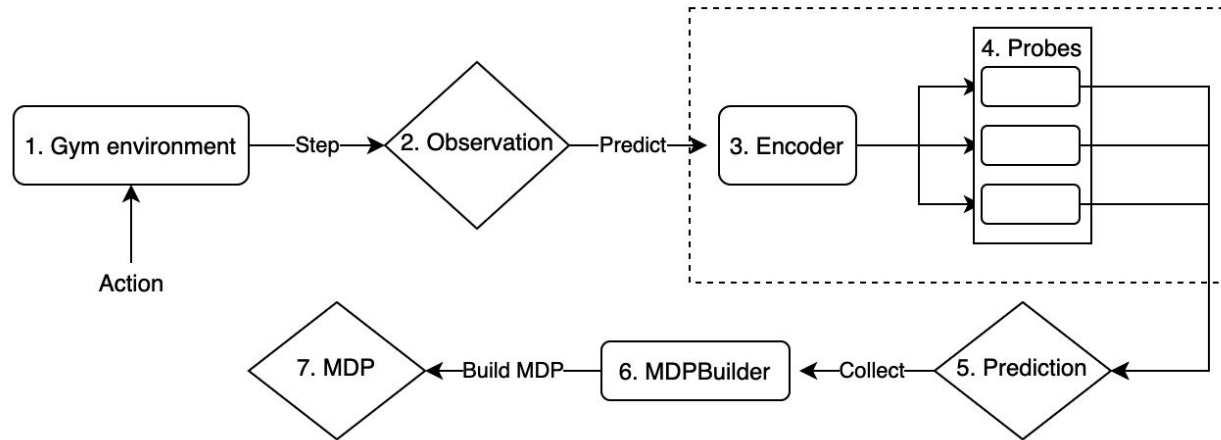
- Encoder/probe training
- Evaluation

Codebase is extended in this project for:

- Obtaining state label predictions
- Building an MDP

Code at: [github.com/davidkerkkamp/representation-learning](https://github.com/davidkerkkamp/representation-learning)

# Pipeline







# Using the encoder

- Code contains class `AtariARIHandler` for:
  - Encoder training
  - Obtaining predictions
- Use handler to obtain Gym environment

```
args.env_name = 'PongNoFrameskip-v4'  
handler = AtariARIHandler(args, wandb)  
gym_env = handler.get_gym_env()
```

Example: create handler and obtain env for game Pong



# Stepping through the Gym environment

- Apply **step** on environment using some **action**
- Each step returns 4 objects:
  - **obs**: 210 x 160 observation matrix representing game screen
  - **reward**
  - **done**: boolean indicating end of episode
  - **info**: dict with ground truth labels of observation

```
for i in range(n):  
    obs, reward, done, info = gym_env.step(action)
```

Example: perform `n` steps on environment  
using some `action`



# Obtaining state predictions

- Handler contains `predict` function
- Observation is put through encoder and probes
- Result is a dictionary with labels

```
prediction = handler.predict(obs)
```

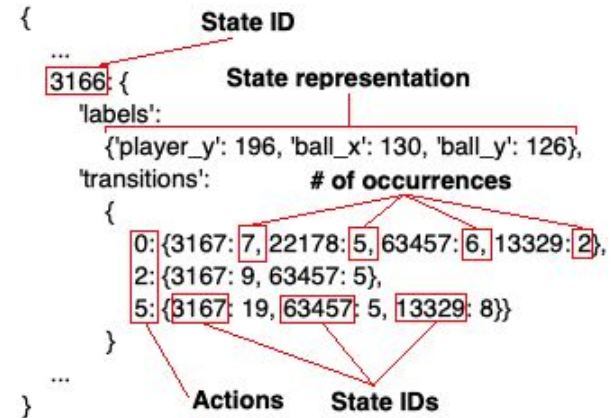
Example: obtain predicted state labels for observation

```
{  
  'player_y': 108,  
  'enemy_y': 196,  
  'ball_x': 156,  
  'ball_y': 55,  
  'enemy_score': 5,  
  'player_score': 0  
}
```

Example of a dictionary with predicted labels

# Building an MDP

- Class `MDPBuilder` for:
  - Collecting observations
  - Building a Markov Decision Process
- Keeps track of:
  - Distinct observations as states
  - Transitions for all actions
- Export MDP to PRISM file
  - Can be read by model checkers PRISM or Storm



Internal data structure of `MDPBuilder` class



# Building an MDP

```
labels_to_use = ['ball_x', 'ball_y', 'player_y']
actions_to_use = [0, 2, 5]
mdp_builder = MDPBuilder(labels_to_use, actions_to_use)

for i in range(n):
    ....
    mdp_builder.add_state_info(prediction, action)
```

Example: create MDPBuilder instance and add observations

**// Command describing transitions for some action**

```
[] player_y=186 & ball_x=205 & ball_y=0 ->
  0.75 : (player_y'=186) & (ball_x'=205) & (ball_y'=0) +
  0.25 : (player_y'=195) & (ball_x'=205) & (ball_y'=0);
```

**// Transitions from same state as above, different action**

```
[] player_y=186 & ball_x=205 & ball_y=0 ->
  0.33 : (player_y'=186) & (ball_x'=205) & (ball_y'=0) +
  0.33 : (player_y'=187) & (ball_x'=205) & (ball_y'=0) +
  0.17 : (player_y'=181) & (ball_x'=205) & (ball_y'=0) +
  0.17 : (player_y'=190) & (ball_x'=205) & (ball_y'=0);
```

**// Different state**

```
[] player_y=187 & ball_x=205 & ball_y=0 ->
  0.33 : (player_y'=192) & (ball_x'=205) & (ball_y'=0) +
  0.33 : (player_y'=187) & (ball_x'=205) & (ball_y'=0) +
  0.33 : (player_y'=184) & (ball_x'=205) & (ball_y'=0);
```

Example of generated PRISM file

# Pong


Action space: {no action, up, down}

Available labels:

- player\_y, enemy\_y
- ball\_x, ball\_y
- player\_score, enemy\_score
- player\_x and enemy\_x are **constants**

Undesired states:

- $\text{ball\_x} > \text{player\_x}$  (enemy scores a point)





# Demo



# Conclusion

Result:

- Pipeline that takes game observations and creates model of environment
- Model can be used to avoid dangerous situations for the player

Future:

- Other MDP model formats (e.g. explicit)
- Way of collecting observations (now random)
- Better ways to generate MDP (similar transitions, small probabilities, states with 1 transition)
- Applying model checking on generated models





# Questions